

Privacy Choices

A privacy decision aid that asks what the student values **before** a service decides for them.

The problem I chose

Privacy policies disclose practices, but the work of interpreting them is left to the student. I wanted to connect those disclosures to five questions: what am I sharing, what am I unwilling to share, how is it handled, what can I change, and do I still want to use this service?

01 Choose

Select a service and local preferences.

02 Understand

Inspect conflicts and source evidence.

03 Reduce

Apply a control or review a guide.

04 Decide

Acknowledge what remains, or decline.

What I built

Privacy Choices Preferences saved locally

YOUR PRIVACY, YOUR CALL

Before you visit

Decide what you are comfortable sharing. Desktop browser only

1 Preferences 2 Findings 3 Reductions 4 Your decision

Choose a service

3 supported services

Service name or URL

Google Maps [Analyze policy →](#)

Google Maps Location Quizlet Tracking ChatGPT Data use Google Maps selected. No company page has been opened.

An access token is required for live analysis. [Configure access token](#)

Only the service ID goes to the analysis backend. Your preferences stay on this device.

Your preferences

LOCAL ONLY

Selected means acceptable to you, not permission granted to the service.

Information I can share

0 of 8 selected [None](#) [All](#)

- ☐ Identity & contact details
Name, email, account details
- ☐ Browser location
Location requested from this browser
- ☐ Approximate location
Area inferred from your IP address
- ☐ Device & cookie identifiers
Browser, device, and session identifiers

POLICY REVIEW

Not requested

No findings yet
No policy analysis has been received for this review.

Selected service	Google Maps
Acceptable information	0 of 8 types
Collection assessment	Not assessed

Actual extension interface. Preferences stay in the browser. Requests send a supported service ID and private access token, not the original URL or preference selections.

Three services, three different privacy decisions

Google Maps

Block browser location; IP location and typed places remain.

Quizlet

Review optional tracking choices; do not promise to remove essential processing.

ChatGPT

Distinguish receiving a conversation from using it for model training.

The decisions behind it

Agency over a score	No universal safe/unsafe rating. The student chooses acceptable data and uses, then acknowledges unresolved risks.
Separate facts from protection	Company policy, AI interpretation, user-reported changes, and browser-verified controls are different states.
Small sections, checked evidence	Cover each section, copy intact evidence, then run a separate AI check. Unsupported details become unknown, not confident guesses.
Small and inspectable	Native extension UI, a small Node backend, and persistent budget reservations. No browsing-history access or autonomous account edits.

Verified behavior, not a blanket privacy guarantee

CHATGPT

Policy findings

Mixed sources

Analyzed Sep 11, 2026, 5:51 AM. US consumer policies, English.

AI interpretations of company disclosures, not an independent audit. Check the evidence; missing information is not proof of no collection.

Company disclosures are not an independent audit or verification of your account settings.

All 5 sections of the available text completed extraction and a separate AI evidence check. Findings remain selective, not a complete inventory.

10 candidate findings were excluded after evidence review or deduplication.

The evidence check is another AI judgment, not independent verification. Categories, explanations, purposes, and retention interpretations may still be wrong.

Public-reader retrieval was used. Upstream completeness and freshness are not independently guaranteed.

Some evidence is a dated snapshot. Its currentness could not be verified.

Some or all evidence comes from dated snapshots. Policy wording and account controls may have changed since capture.

1 conflicts

0 browser-verified

0 within preferences

Content used to train models

Preference conflict

User-submitted content may be used to train OpenAI models.

Text & uploaded files is not selected as acceptable. Unaccepted uses: AI model training.

► Details & policy evidence

Actual hosted result. The ChatGPT view connects a training disclosure to unaccepted content and model-training preferences. Source and uncertainty warnings stay visible.

57 automated checks passed. Includes an isolated real-Redis test: 300 concurrent small reservations accepted 250, stopping at \$5. Earlier spending is preserved.

Evidence boundary. This capture uses a live hosted response, not synthetic findings. Browser-location controls are tested separately with labeled fixtures. Neither establishes actual company behavior.

Feasibility and return on effort

Publicly accessible: policies and help pages about collection, purposes, retention, and controls. Five dated source snapshots were obtained. Latest retrieval used two direct sources, one reader, and two snapshots. **Not established:** internal processing, complete tracking prevention, account-specific choices, or semantic correctness of AI summaries.

Return: three bounded services make the preference-to-decision flow testable. Broader coverage would add recurring policy review and site-specific maintenance. Live model calls are now exercised; each extraction and evidence check reserves budget against the original \$5 allowance. Educational usefulness is a hypothesis: no student usability study has been completed.

Current delivery status: extension ZIP built; backend deployed and authentication checked. Hosted checks returned source-linked findings for all three services. AI interpretations are selective and fallible; actual collection and account settings remain unverified.

Sources and evidence: Google privacy / location | Quizlet privacy | ChatGPT data controls | Chrome content settings

Public policies were captured September 11, 2026 (UTC). Source details, test commands, and the remaining setup steps accompany the code in README.md and DECISIONS.md.